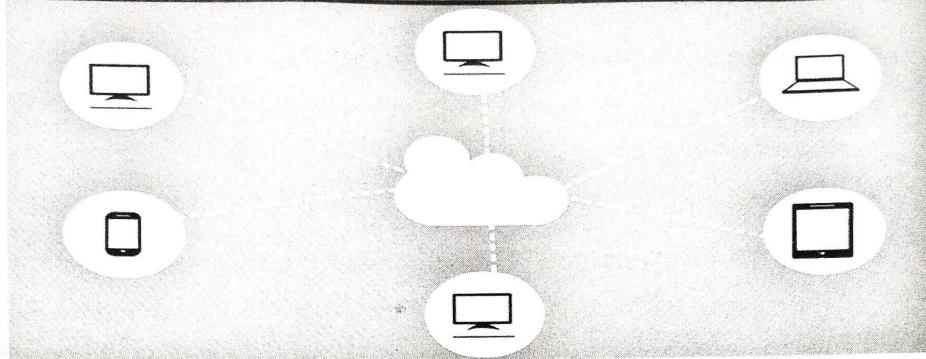


Chapter

3



INTER-PROCESS COMMUNICATION

3.1 INTRODUCTION

The process of exchanging the data between two or more independent process in a distributed environment is called as Inter Process Communication (IPC). In a distributed system, communication between processes is crucial for sharing information, coordinating actions, and ensuring overall system functionality. Processes in a distributed system are units of execution that can run on different nodes (computers) connected via a network.

Mechanisms of Inter Process Communications

Effective communication between these processes involves several key concepts and mechanisms:

- **Message Passing:** Processes communicate by sending messages to each other. Message passing can be synchronous or asynchronous. In synchronous communication, the sender waits for a response before proceeding, while in asynchronous communication, the sender continues its execution without waiting for a reply.

- **Remote Procedure Calls (RPC):** RPC is a protocol that allows a process to execute a procedure (function) on a remote node as if it were a local call. The caller sends a message containing the necessary parameters to the remote process, which executes the procedure and returns the result.
- **Message-Oriented Middleware (MOM):** MOM is a software layer that facilitates communication between distributed processes by managing the sending, receiving, and queuing of messages. It provides a way for processes to exchange messages without worrying about the underlying network complexities.
- **Synchronization and Coordination:** Distributed systems require mechanisms for synchronizing activities among processes to ensure consistency and avoid conflicts. Techniques like locks, semaphores, and distributed algorithms (e.g., distributed mutual exclusion, distributed consensus) help in coordinating actions across multiple processes.
- **Communication Protocols:** Various communication protocols are used in distributed systems, such as TCP/IP, UDP, HTTP.
- **Data Serialization:** When messages are sent between processes, the data needs to be serialized (converted into a format that can be transmitted over the network) and deserialized at the receiving end. Common serialization formats include JSON, XML, Protocol Buffers, and MessagePack.
- **Fault Tolerance and Reliability:** Distributed systems must handle failures gracefully. Techniques like replication, fault detection, and recovery mechanisms are employed to ensure reliability and fault tolerance in communication.
- **Scalability and Performance:** Communication protocols and mechanisms must be designed to scale efficiently as the system grows. Load balancing and optimizing message routing contribute to better performance in distributed communication.
- **Security:** Secure communication between processes involves encryption, authentication, and authorization mechanisms to protect data integrity, confidentiality, and prevent unauthorized access.
- **Discovery and Addressing:** Processes in a distributed system need to discover each other and communicate. Service discovery mechanisms and addressing schemes (such as IP addresses, ports, or domain names) help in locating and reaching the intended processes.

In summary, communication between processes in a distributed system involves various complex mechanisms and considerations to ensure effective, reliable, and secure exchange of information across networked nodes.

Benefits of Inter Process Communications

Communication between processes in a distributed system offers numerous benefits that are fundamental to the functionality, efficiency, and reliability of the system.

Some of the key benefits include:

- **Resource Sharing:** Processes in a distributed system can share resources such as data, computation, and services by communicating with each other. This enables efficient utilization of resources across the network.
- **Improved Performance:** Communication allows distributed processes to collaborate on tasks, leading to improved performance by leveraging distributed computing power. Parallel processing and workload distribution enhance overall system performance.
- **Scalability:** Distributed systems can easily scale by adding or removing nodes or processes. Communication enables seamless integration of new resources and accommodates increased workloads without significant architectural changes.
- **Fault Tolerance:** Through communication, distributed systems can replicate data and services across multiple nodes. This redundancy helps in achieving fault tolerance and high availability. If a node fails, other nodes can continue functioning, ensuring system reliability.
- **Increased Flexibility:** Processes can be distributed across various locations or devices, allowing flexibility in system design and deployment. This flexibility enables the system to adapt to changes in workload or infrastructure without major disruptions.
- **Decentralization:** Communication facilitates decentralization, allowing processes to operate independently while cooperating to achieve common goals. This decentralization leads to better load distribution and reduces the reliance on a single point of failure.

- Global Reach:** Distributed systems enable communication between processes across geographical locations. This global reach facilitates collaboration among users and systems regardless of their physical locations.
- Reduced Latency:** By placing processes closer to the data or services they require, distributed systems can reduce communication latency. This can result in faster response times and better user experiences.
- Cost Efficiency:** Distributed systems often leverage existing resources efficiently without the need for large, centralized infrastructure. This cost-effective approach allows for better resource utilization and reduces operational expenses.
- Support for Concurrent Operations:** Communication between distributed processes allows for concurrent execution of tasks, enabling parallelism and better handling of multiple requests simultaneously.

3.2 THE API FOR THE INTERNET PROTOCOLS

In this section, we discuss the general characteristics of inter-process communication and then discuss the Internet protocols. The Internet is built upon various protocols that enable communication and data transfer between devices. Regarding APIs for these protocols, there isn't a single API for the entire suite of Internet protocols. However, there are APIs and libraries available that developers can use to interact with these protocols programmatically.

3.2.1 Characteristics of Inter-Process Communication

There are mainly five characteristics of inter-process communication in a distributed environment/system.

1. Synchronous Communication:

- In synchronous communication, the sender process waits until the receiver acknowledges receipt of the message or completes the requested operation before continuing its execution.
- It involves a blocking mechanism where the sender is paused until a response is received, which can potentially introduce delays if the receiver is slow or unresponsive.

- This method ensures a predictable order of operations and can simplify error handling since the sender knows when the message has been processed.
- Asynchronous Communication:**
 - Asynchronous communication allows the sender to continue its execution immediately after sending a message without waiting for a response from the receiver.
 - It involves a non-blocking mechanism, enabling processes to operate independently, enhancing concurrency and potentially improving system performance.
 - Asynchronous communication requires specific protocols or mechanisms to handle responses or acknowledgment, such as callbacks, notifications, or polling.

3. Message Destination

- Ports serve as gateways for messages to enter or exit processes. Each port has a unique identifier, allowing processes to send or receive messages through these designated channels. When a process initiates communication, it specifies the target port where the message should be delivered. The receiving process monitors its designated port, enabling the directed reception of messages. By addressing messages to specific ports, IPC ensures precise routing and delivery, facilitating organized and efficient communication among processes in distributed systems, bolstering reliability, and supporting various interactions within the network.
- A local port is a message destination within a computer, specified as an integer. A port has exactly one receiver but many senders. Processes may use multiple ports from which to receive messages. Any process that knows the number of a port can send the message to it.

4. Reliability

- Reliability in Inter-process communication (IPC) ensures accurate and consistent message delivery between processes.
- IPC mechanisms incorporate error detection, retransmission, and acknowledgment protocols to guarantee message integrity and successful

transmission. These features help mitigate data loss, ensure ordered delivery, and maintain communication reliability, vital for dependable and fault-tolerant operations within distributed systems.

5. Ordering

- It is the process of delivering messages to the receiver in a particular order. Some applications require messages to be delivered in the sender order i.e. the order in which they were transmitted by the sender.
- IPC mechanisms enforce message ordering through techniques like sequence numbers or timestamps. Preserving message sequence ensures that received messages reflect the same order as they were sent, crucial for consistent data processing and accurate synchronization among processes in distributed systems.

3.2.2 Sockets

Sockets are a fundamental mechanism for Inter-process communication (IPC) used for establishing communication channels between processes, either on the same machine or across a network. They facilitate bidirectional data flow between processes by creating endpoints for sending and receiving data. Sockets utilize network protocols (such as TCP/IP or UDP) and provide APIs for processes to establish connections, exchange messages, and close connections when communication is complete.

Here's a detailed overview of sockets in IPC:

1. Types of Sockets:

- Stream Sockets (TCP):** Reliable, connection-oriented sockets using Transmission Control Protocol (TCP). They ensure ordered and error-checked delivery of a stream of bytes.
- Datagram Sockets (UDP):** Unreliable, connectionless sockets using User Datagram Protocol (UDP). They allow sending discrete messages (datagrams) without guaranteed delivery or order.

2. Socket APIs:

- Programming interfaces (APIs) like BSD sockets, Winsock (Windows), and socket-related libraries in various programming languages (e.g., Python's

`socket` module) provide methods to create, bind, listen, accept, connect, send, and receive data via sockets.

3. Socket Operations:

- Creation:** Sockets are created using system calls (e.g., `socket()` in C).
- Binding:** Processes associate a socket with a specific address and port using `bind()` to listen for incoming connections or specify the source of outgoing connections.
- Listening and Accepting:** Servers call `listen()` to await incoming connections and `accept()` to establish connections with clients, creating a new socket for communication.
- Connecting:** Clients initiate connections to servers using `connect()` to establish communication.
- Sending and Receiving:** Processes use `send()` and `recv()` (or equivalent functions) to transfer data over sockets. For UDP, `sendto()` and `recvfrom()` are used due to the connectionless nature.

4. Socket Communication Models:

- Client-Server:** Commonly used model where servers passively listen for incoming connections, and clients actively initiate connections.
- Peer-to-Peer:** Processes communicate directly without a central server, often seen in decentralized applications.

5. Socket Characteristics:

- Sockets possess attributes such as address, port number, protocol, and communication characteristics (e.g., TCP's reliability or UDP's simplicity and lower overhead).

3.2.3 UDP Datagram Communication

UDP is connectionless and unreliable internet protocol mainly used to send short messages called *datagrams*. It doesn't require making a connection with the host to exchange data. Since UDP is unreliable protocol, there is no mechanism for ensuring that data sent is received.

UDP Features

UDP communication involves the following aspects:

- **Connectionless Communication:** Unlike TCP, UDP is connectionless. It doesn't establish a connection before sending data; instead, it directly sends datagrams to the recipient without prior negotiation.
- **Datagram Structure:** UDP transmits data in datagrams. Each datagram consists of a header (containing source and destination port numbers, length, and a checksum for error detection) and the payload (the actual data being sent).
- **Unreliable Delivery:** UDP does not guarantee delivery, ordering, or acknowledgment of packets. There's no mechanism for retransmission or ensuring that all data arrives in order at the destination. This makes UDP faster but less reliable compared to TCP.
- **Low Overhead:** Because UDP lacks the overhead associated with managing connections and ensuring reliability, it offers lower latency and is more lightweight, making it suitable for applications where speed is prioritized over reliability (e.g., real-time multimedia streaming, online gaming).
- **Broadcast and Multicast Support:** UDP allows broadcasting messages to multiple recipients (broadcast) or specific groups of recipients (multicast).

Issues of UDP

Datagram communication, particularly with protocols like UDP (User Datagram Protocol), presents several inherent issues and challenges:

- **Unreliable Delivery:** Datagram communication lacks guaranteed delivery. Packets can be lost, duplicated, or delivered out of order without notification, making it unsuitable for applications requiring reliable, ordered data transmission.
- **No Congestion Control:** UDP doesn't incorporate congestion control mechanisms. Consequently, it might lead to network congestion if data is sent at a rate higher than the network capacity, potentially resulting in packet loss.
- **Limited Packet Size:** Datagram protocols like UDP have a limited packet size (65,535 bytes), and if the data exceeds this limit, fragmentation might occur, which can impact network efficiency.

- **No Flow Control:** There's no inherent flow control mechanism in UDP. This means that a fast sender can overwhelm a slower receiver, leading to data loss or buffer overflow issues at the receiver's end.
- **Security Concerns:** Datagram communication lacks built-in encryption or acknowledgment mechanisms, making it susceptible to security threats like eavesdropping, data tampering, or spoofing. Extra security measures need to be implemented at the application level.
- **Handling of Error Conditions:** Datagram communication doesn't inherently support error handling. If errors occur during transmission, there's no automatic recovery mechanism, and applications need to handle errors themselves.
- **Routing and Delivery Issues:** Datagram packets might take different routes to reach the destination, which can lead to varying delivery times (packet jitter). This unpredictability can be a concern for real-time applications.
- **Application-Level Overhead:** Due to its unreliable nature, applications using datagram communication (like UDP) need to implement their own error checking, retransmission, and sequencing mechanisms, adding complexity to application code.

3.2.4 TCP Stream Communication

TCP (Transmission Control Protocol) is a connection-oriented communication protocol that operates at the transport layer of the TCP/IP model. TCP provides reliable, ordered, and stream-oriented communication between two hosts.

Features of TCP Stream Communication

Here are the key aspects of TCP stream communication:

- **Connection-Oriented:** TCP establishes a connection between the sender and receiver before transmitting data. It follows a three-way handshake (SYN, SYN-ACK, ACK) to establish and terminate connections.
- **Reliable Delivery:** TCP guarantees reliable delivery of data. It ensures that data sent by one party is received accurately and in the same order by the other party. It employs acknowledgments, sequence numbers, and retransmissions to achieve reliability.
- **Ordered Delivery:** TCP preserves the order of data transmission. Data sent in

multiple segments from one end is received in the same order at the other end, preventing issues like packet reordering or data corruption.

- **Stream-Oriented Communication:** TCP sends data in a continuous stream of bytes. The sender divides data into smaller segments and numbers them for transmission. At the receiver's end, TCP reassembles these segments back into the original data stream.
- **Flow Control:** TCP implements flow control mechanisms to prevent overwhelming the receiver. It uses a sliding window approach, allowing the sender to transmit a limited number of unacknowledged segments based on the receiver's buffer capacity.
- **Congestion Control:** TCP manages network congestion by adapting the transmission rate to avoid overwhelming the network. It uses algorithms to regulate the transmission rate based on network conditions, ensuring fair and efficient data transfer.
- **Error Checking and Retransmission:** TCP performs error checking using checksums to detect corrupted data. If errors are detected or acknowledgments are not received within a specified time, TCP triggers retransmission of the data.
- **Full Duplex Communication:** TCP supports full-duplex communication, allowing simultaneous two-way data flow. Both the sender and receiver can send and receive data independently within the same connection.

Issues of TCP

TCP stream communication, while robust and reliable, also presents certain challenges and issues:

- **Latency:** TCP's connection setup, acknowledgments, and flow control mechanisms can introduce latency, impacting real-time or delay-sensitive applications.
- **Overhead:** TCP carries additional overhead due to features like sequence numbers, acknowledgments, and error-checking mechanisms, which might affect performance, especially for small data transfers.
- **Head-of-Line Blocking:** If a packet within a TCP stream encounters an issue of delay, subsequent packets must wait, causing head-of-line blocking and potentially affecting the overall throughput.

- **Congestion Control:** TCP's congestion control mechanisms may sometimes be overly conservative, leading to reduced data transmission rates in situations where the network could handle more traffic.
- **Buffering and Buffer Overflow:** TCP uses buffers to manage incoming data. In high-throughput scenarios, buffer overflow might occur if data arrives faster than it can be processed, potentially resulting in packet loss.
- **Suitability for Real-Time Applications:** TCP's focus on reliability and ordered delivery may not be ideal for real-time applications like video streaming or gaming due to added latency and strict packet ordering.

API for TCP Stream Communication

TCP stream communication is commonly facilitated through various programming language-specific APIs or libraries that provide functions or classes to create TCP connections, send and receive data, manage sockets, and handle TCP-specific functionalities. Some popular APIs and libraries for TCP stream communication include:

- **C/C++ Sockets API (BSD Sockets):** The standard sockets API in C and C++ provides functions like `socket()`, `bind()`, `listen()`, `accept()`, `connect()`, `send()`, and `recv()` to create, manage, and communicate over TCP sockets.
- **Java Socket API:** Java provides the `java.net` package that includes classes like `Socket` and `ServerSocket` for TCP communication. These classes allow creating client and server sockets, establishing connections, and exchanging data using input/output streams (`InputStream` and `OutputStream`).
- **Python socket module:** Python's socket module enables TCP socket programming with functions like `socket()`, `bind()`, `listen()`, `accept()`, `connect()`, `send()`, and `recv()`. It allows developers to create TCP sockets, establish connections, and perform data transfer.
- **C# .NET Socket Class:** The .NET framework offers the `System.Net.Sockets`. `Socket` class for TCP communication. This class provides methods to create, connect, send, and receive data over TCP sockets in C#.
- **Node.js Net Module:** Node.js provides the `net` module for TCP communication. Developers can create TCP servers and clients using functions like `net.createServer()` and `net.connect()` to handle connections and data transfer.

3.3 EXTERNAL DATA REPRESENTATION

External Data Representation (EDR) is a method used to represent data in a standardized format that allows it to be exchanged between different systems, platforms, or programming languages. It provides a common way to serialize data (convert it into a format suitable for transmission) and deserialize it (reconstruct it back into its original form) regardless of the underlying architecture or environment.

Computers handle a variety of data types, and these types vary depending on the location and purpose of the data transmission. Individual primitive data items can have a variety of data values, and not all computers store primitive values like integers in the same order. Floating-point numbers are also represented differently by different architectures. There are two ways to arrange integers: big-endian order, which puts the Most Significant Byte (MSB) first; and little-endian order, which either puts the MSB last or the Least Significant Byte (LSB) first. The collection of codes used to represent characters is still another problem. Most applications on UNIX systems use ASCII character coding, which uses one byte per character, whereas the Unicode standard uses two bytes per character and allows for the representation of texts in many different languages.

One of the following methods can be used to enable any two computers to exchange binary data values:

- The values are converted to an agreed external format before transmission and converted to the local form on receipt; if the two computers are known to be the same type, the conversion to external format can be omitted.
- The values are transmitted in the sender's format, together with an indication of the format used, and the recipient converts the values if necessary.

3.3.1 Characteristics of External Data Representation:

- Platform-Independence:** EDR ensures that data can be represented in a way that is independent of the hardware, operating system, or programming language used. This enables data interchangeability across diverse systems.
- Serialization:** EDR formats enable the serialization of complex data structures into a format suitable for transmission or storage. This serialized form can be easily reconstructed at the receiving end.

- Standardized Format:** EDR typically follows a standardized format or schema, which facilitates both readability for humans and parsing for machines.
- Efficiency:** EDR formats aim to strike a balance between human readability (in some cases) and efficiency in terms of size and processing speed. For instance, some formats like JSON and XML are human-readable but may not be as space-efficient as binary formats like Protocol Buffers or MessagePack.

Marshalling and Unmarshalling

Marshalling, also known as serialization, is the process of transforming data structures or objects from their in-memory representation into a format that is suitable for transmission or storage. This involves converting complex data types or objects into a simpler and standardized format, typically to be transmitted over a network or stored persistently. During marshalling, the data is encoded into a specific format, such as JSON, XML, Protocol Buffers, or binary representations, facilitating its transfer across different systems or platforms.

On the other hand, unmarshalling, also referred to as deserialization, is the reverse process. It involves reconstructing the serialized data back into its original in-memory representation or data structure. Unmarshalling takes the encoded data obtained from transmission or storage, decodes it into a format that the system can understand, and then reconstructs the complex data types or objects from the serialized form.

Together, marshalling and unmarshalling ensure the seamless transmission, storage, and reconstruction of data between heterogeneous systems or applications. These processes are fundamental in enabling interoperability, as they allow different systems, developed using various languages or running on different platforms, to communicate and exchange data by converting it into a universally understandable format and back again.

3.4 EDR APPROACHES

There are three ways to successfully communicate between various sorts of data between computers.

3.4.1 CORBA's Common Data Representation (CDR)

CORBA (Common Object Request Broker Architecture) uses a standardized data representation called Common Data Representation (CDR) to facilitate the exchange of data between different computing environments in a distributed system.

CDR serves as a protocol-neutral binary encoding format used by CORBA for serializing and deserializing data, allowing different programming languages and platforms to communicate seamlessly. It defines rules for representing various data types in a way that ensures compatibility and interoperability between systems.

Key features of CORBA's Common Data Representation (CDR):

- **Data Types Mapping:** CDR defines rules for mapping different data types used in various programming languages (such as integers, floats, strings, arrays, structures, etc.) into a standardized binary format. It ensures that these data types are represented consistently across different systems.
- **Serialization (Marshalling):** When data needs to be transmitted or stored, CDR serializes the data by encoding it into a binary format. The process involves converting the in-memory data structures or objects into a stream of bytes adhering to the CDR format. This serialized data can then be transmitted across the network or stored persistently.
- **Binary Encoding:** CDR employs a binary encoding scheme optimized for efficiency in transmission and storage within a distributed system. It encodes data in a way that accounts for differences in byte order (endianness) between systems, ensuring proper interpretation at the receiving end.
- **Transmission and Reception:** Serialized data encoded using CDR is transmitted across the network between different CORBA components or systems. At the receiving end, the transmitted data is received and subjected to the deserialization process.
- **Deserialization (Unmarshalling):** Upon receiving the serialized data, CDR performs deserialization by decoding the binary stream back into its original data structure or object format. This involves reconstructing the in-memory representation of the data, adhering to the CDR rules for data type mapping and byte order, to ensure accuracy and consistency.
- **Interoperability:** CDR's standardized binary representation enables CORBA-based applications implemented in different languages and running on diverse platforms to communicate effectively. It ensures that data can be serialized, transmitted across heterogeneous systems, and accurately reconstructed at the receiving end, fostering interoperability.

In essence, CORBA's Common Data Representation (CDR) serves as a crucial mechanism for encoding, transmitting, and decoding data in a standardized binary format within a distributed CORBA environment, enabling seamless communication and interoperability between heterogeneous systems.

3.4.2 Java Object Serialization

In Java, object serialization serves as a means to convert Java objects into a stream of bytes for storage or transmission. This serialization mechanism acts as an external data representation, allowing objects to be stored persistently or transmitted across a network in a platform-independent format.

When you use Java's object serialization, the serialized object is represented as a stream of bytes. This serialized form can be saved into a file, sent over a network, or stored in a database. The key advantages of using Java object serialization as an external data representation include:

- **Platform Independence:** Serialized objects can be written and read across different platforms since the serialization format isn't tied to a specific system or architecture.
- **Persistence:** Serialized objects can be stored persistently in files or databases and reconstructed later as needed, retaining their state and structure.
- **Data Transmission:** Serialized objects can be sent across a network or between different applications as a stream of bytes, enabling communication between disparate systems.

Here's a breakdown of how Java object serialization works as an external data representation:

1. **Serialization:** During serialization, Java inspects the object's class metadata and converts its state (fields and their values) into a sequence of bytes. Transient fields and non-serializable objects are excluded or handled separately.
 - (a) **Implementing Serializable:** For a Java class to be serializable, it needs to implement the `java.io.Serializable` interface. This is a marker interface that indicates to the Java Virtual Machine (JVM) that objects of this class can be serialized.

```

import java.io.Serializable;

public class MyClass implements Serializable
{
    // class members, constructors, methods
}

```

(b) **Serialization Process:**

- Create an instance of **ObjectOutputStream**, which writes Java objects into a byte stream.
- Use **ObjectOutputStream** to serialize the object into a sequence of bytes.
- The object's state (fields and their values) is converted into a platform-independent byte stream.

```

try {
    MyClass obj = new MyClass(); // Object to be serialized
    FileOutputStream fileOut = new FileOutputStream("file.ser");
    ObjectOutputStream out = new ObjectOutputStream(fileOut);
    out.writeObject(obj); // Serialize the object
    out.close();
    fileOut.close();
    System.out.println("Object serialized successfully");
} catch (IOException e) {
    e.printStackTrace();
}

```

2. Deserialization: During deserialization, Java reconstructs the object by reading the serialized byte stream. It recreates the object's state and structure based on the serialized data, initializing the object's fields with the deserialized values.

(a) **Deserialization Process:**

- Create an instance of **ObjectInputStream**, which reads serialized Java objects from a byte stream.
- Use **ObjectInputStream** to deserialize the byte stream back into an object.
- Reconstruct the object by restoring its state from the byte stream.

```

try {
    FileInputStream fileIn = new FileInputStream("file.ser");
    ObjectInputStream in = new ObjectInputStream(fileIn);
}

```

```

MyClass obj = (MyClass) in.readObject(); // Deserialize the
object
in.close();
fileIn.close();
System.out.println("Object serialized successfully");
} catch (IOException | ClassNotFoundException e) {
    e.printStackTrace();
}

```

3.4.3 XML (Extensible Markup Language)

XML (Extensible Markup Language) serves as a flexible and standardized format for representing structured data. When used in external data presentation, XML works by providing a set of rules for encoding documents in a format that is both human-readable and machine-readable. Here's how it typically works:

1. **Syntax and Structure:** XML uses tags to define elements and attributes that structure the data. Tags are enclosed in angle brackets, and they represent the hierarchy and relationships between different pieces of information. For instance:

```

<Catalog>
    <Book>
        <Title>CSA</Title>
        <Author>S. Bhardwaj</Author>
    </Book>
    <Book>
        <Title>E-Commerce</Title>
        <Author>P. Kumar</Author>
    </Book>
</Catalog>

```

In this example, **<Catalog>** contains multiple **<book>** elements, each with **<title>**, and **<author>** sub-elements.

2. **Well-Formed and Validity:** XML documents must adhere to specific rules to be considered well-formed. They should have a single root element, properly nested tags, and quoted attribute values. Additionally, XML can have a defined structure through Document Type Definitions (DTDs) or XML Schemas, ensuring document validity against a predefined structure.

3. **Data Exchange and Integration:** XML facilitates data exchange between different systems or platforms. When one system needs to share data with another, it can serialize its data into XML format. The recipient system can then parse the XML and extract the relevant information. This interoperability is crucial for integrating diverse systems that need to communicate and share information.
4. **Parsing and Processing:** XML parsers are used to read, interpret, and process XML documents. These parsers can be part of programming languages or libraries, and they transform the XML data into a format that can be manipulated within the application or system. For instance, in web development, JavaScript can parse XML using the DOM (Document Object Model) to access and modify XML data.
5. **Presentation and Display:** External presentation of XML data can take various forms. It can be transformed into other formats like HTML using XSLT (Extensible Stylesheet Language Transformations) or styled using CSS (Cascading Style Sheets) for web display. Additionally, XML data can be converted to JSON, CSV, or other formats based on the requirements of the presentation layer.
6. **Web Services and APIs:** XML is widely used in defining APIs and web services. For instance, SOAP (Simple Object Access Protocol) uses XML for structuring messages exchanged between web services.

3.5 MULTICAST COMMUNICATION

Multicast communication refers to a communication paradigm where data is sent from one sender to multiple receivers in a network simultaneously. Unlike unicast (one-to-one) or broadcast (one-to-all) communication, multicast allows for efficient transmission to a specific group of recipients interested in receiving the data.

Here's an overview of multicast communication:

I. Group-based Transmission:

- **Sender:** There is a single sender, also known as the source, which transmits data to a designated group of receivers.
- **Receivers:** Multiple receivers express interest in receiving the multicast data by joining a specific multicast group.

2. IP Multicast:

- *IP Multicast Addressing:* IP multicast uses special addresses (in the range of 224.0.0.0 to 239.255.255.255) to define multicast groups. Data sent to these addresses is distributed to all members of the corresponding multicast group.
- *Efficiency:* It conserves network bandwidth as data is transmitted only once by the sender, and routers replicate and forward the data to multicast group members as needed.

3. Applications of Multicast Communication:

- *Streaming Media:* Live audio/video streaming applications often use multicast to deliver content to multiple viewers simultaneously.
- *Financial Services:* Stock market data distribution to subscribers, where many subscribers are interested in real-time updates, can benefit from multicast communication.
- *Content Distribution Networks (CDNs):* CDNs can use multicast to efficiently distribute content to various nodes across the network.

4. Protocols and Technologies:

- *IGMP (Internet Group Management Protocol):* Used by hosts to join or leave multicast groups, allowing routers to manage multicast group membership.
- *PIM (Protocol Independent Multicast):* Routing protocols that support multicast communication by efficiently forwarding multicast packets across the network.

5. Challenges:

- *Scalability:* Managing large multicast groups might pose scalability challenges for networks and systems.
- *Security:* Ensuring secure communication within multicast groups can be more complex compared to unicast communication.

6. Use in Various Network Layers:

- *Transport Layer:* Multicast can be implemented in transport layer protocols like UDP (User Datagram Protocol).

- *Network Layer:* IP multicast operates at the network layer and relies on routers' ability to handle multicast traffic.

7. Adoption and Support:

- *Internet Service Providers (ISPs):* Some ISPs support multicast for specific applications or services.
- *Enterprise Networks:* Multicast is used in various enterprise settings, especially for applications that require efficient group communication.

Multicast communication offers an efficient way to transmit data to multiple recipients across networks, reducing network load and duplicative transmissions. However, its adoption might be limited by network infrastructure, security considerations, and application-specific requirements.

3.6 NETWORK VIRTUALIZATION: OVERLAY NETWORKS

Network Virtualization is a process of logically grouping physical networks and making them operate as single or multiple independent networks called Virtual Networks.

Network virtualization software allows network administrators to move virtual machines across different domains without reconfiguring the network. The software creates a network overlay that can run separate virtual network layers on top of the same physical network fabric.

Network virtualization enables the virtual provisioning of an entire network and separates network services from the underlying hardware. It makes it possible to programmatically create, provision, and manage networks all in software, while continuing to leverage the underlying physical network as the packet-forwarding backplane. Physical network resources, such as switching, routing, firewalling, load balancing, virtual private networks (VPNs), and more, are pooled, delivered in software, and require only Internet Protocol (IP) packet forwarding from the underlying physical network.

In compliance with networking and security policies established for every linked application, network and security services in software are dispersed to a virtual layer (hypervisors, in the data center) and "attached" to specific workloads, such as your virtual machines (VMs) or containers. When a workload is moved to another host, network services and security policies move with it. And when new workloads are created to scale an

application, necessary policies are dynamically applied to these new workloads, providing greater policy consistency and network agility.

Here are some of the key benefits of network virtualization:

- **Resource Optimization:** Efficient use of physical resources by allocating them dynamically based on demand.
- **Isolation and Segmentation:** Virtual networks can be isolated from one another, providing security and segmentation for different applications or user groups.
- **Scalability and Flexibility:** Easy scalability and adaptability to changing network requirements without significant physical changes.
- **Cost-Efficiency:** Reducing hardware costs by maximizing utilization and allowing multiple network environments on the same infrastructure.

3.6.1 Overlay Networks

Network virtualization, particularly through overlay networks, involves the creation of logical network structures that operate independently of the physical network infrastructure. Overlay networks enable the establishment of multiple virtualized networks over a shared physical network without requiring changes to the underlying hardware.

Key Aspects of Overlay Networks in Network Virtualization:

- **Virtual Network Creation:** Overlay networks create virtualized instances that run on top of the physical infrastructure, allowing the creation of multiple logically isolated networks (virtual overlays).
- **Encapsulation:** Traffic within overlay networks is encapsulated, meaning packets from the virtual network are wrapped within the original network's packets. Protocols like VXLAN (Virtual Extensible LAN), GRE (Generic Routing Encapsulation), and others are used for encapsulation.
- **Independence from Underlying Infrastructure:** Virtual networks in overlays operate independently, allowing different configurations, addressing schemes, and policies compared to the physical network.
- **Flexibility and Segmentation:** Overlay networks offer segmentation and isolation, allowing diverse network architectures, different security protocols, and custom routing within each overlay.

- Software-Defined Networking (SDN):** Overlay networks often leverage SDN principles, separating the control plane from the data plane, enabling centralized control and management of virtual network traffic.
- Scalability and Efficiency:** Overlay networks facilitate scalability by allowing the creation of numerous virtual networks on the same physical infrastructure, enhancing resource utilization.

Use Cases and Applications:

- Cloud Computing:** Overlay networks in cloud environments provide virtualized networking for multiple tenants, ensuring network isolation and security.
- Data Center Virtualization:** Overlay networks are used to create isolated environments for different applications or departments within data centers, optimizing resource allocation and management.
- Multi-Tenant Environments:** Service providers leverage overlay networks to offer dedicated virtual networks to different clients or users while utilizing a shared physical network infrastructure.
- Software-Defined Networking (SDN):** SDN implementations often use overlay networks to create programmable, agile, and scalable network architectures.

Overlay networks are a fundamental aspect of network virtualization, offering enhanced flexibility, scalability, and isolation while enabling the coexistence of multiple virtual networks on a shared physical infrastructure. They play a crucial role in modern network architectures, particularly in cloud computing, data centers, and SDN environments.

3.7 CASE STUDY: MPI

The Message Passing Interface (MPI) is a standardized and widely used communication protocol and library specification used in parallel computing and distributed memory architectures. It allows multiple processes or nodes in a distributed computing environment to communicate and exchange data by sending and receiving messages.

The MPI standard defines the user interface and functionality, in terms of syntax and semantics, of a standard core of library routines for a wide range of message-passing capabilities. It defines the logic of the system but is not implementation specific.

The specification can be efficiently implemented on a wide range of computer architectures. It can run on distributed-memory parallel computers, a shared-memory parallel computer a network of workstations, or, indeed, as a set of processes running on a single workstation. The standard has come from a unification of concepts and most attractive ideas from vendor specific message passing system variants through the work of the MPI forum.

Key Features of MPI:

- Parallel Computing:** MPI is designed for parallel computing environments where multiple processes work together to solve a problem or perform computations.
- Message Passing:** Processes communicate through messages, allowing data exchange between nodes in a parallel system.
- Standardized Interface:** MPI provides a set of standard routines and functions that enable programmers to develop parallel applications irrespective of the underlying hardware or architecture.
- Point-to-Point Communication:** MPI supports point-to-point communication between specific processes using functions like `MPI_Send` and `MPI_Recv`.
- Collective Communication:** It offers collective communication operations, allowing synchronization and data exchange between multiple processes, such as broadcasting data to all processes (`MPI_Bcast`), reducing data across processes (`MPI_Reduce`), or gathering data from multiple processes (`MPI_Gather`).
- Support for Various Languages:** MPI is available for multiple programming languages, including C, C++, Fortran, and Python, among others.
- Scalability:** It is designed to scale efficiently to large computing clusters, enabling high-performance computing (HPC) applications.

MPI Implementations:

Various implementations of MPI exist, such as Open MPI, MPICH, Intel MPI, and others. These implementations provide libraries and tools compliant with the MPI standard, offering different optimizations and features suited for different computing environments.

Use Cases:

MPI is commonly used in scientific computing, computational fluid dynamics, weather forecasting, molecular modeling, and other applications that require massive computational power. It allows efficient utilization of distributed computing resources and facilitates the development of parallel algorithms for complex problems.

Overall, MPI serves as a crucial tool for parallel programming, enabling efficient communication and coordination among processes in distributed memory systems, essential for achieving high-performance computing and solving computationally intensive tasks.

QUESTIONS)***Short Answer Questions*****Q1. What is Inter-process Communication (IPC)?**

Ans: Inter-process Communication (IPC) refers to mechanisms and techniques that enable communication and data exchange between different processes running concurrently within an operating system or across a network. IPC facilitates the sharing of information, synchronization, and coordination among these processes, allowing them to communicate, send messages, and share resources efficiently for collaborative tasks.

Q2. Explain the difference between synchronous and asynchronous IPC mechanisms.

Ans: Synchronous IPC involves a sender waiting for the recipient to receive, process, and acknowledge a message before continuing, ensuring real-time interaction but potentially causing delays. Asynchronous IPC allows the sender to proceed without immediate confirmation, enhancing efficiency by permitting concurrent tasks, albeit without immediate response assurance, reducing potential wait times between communication events.

Q3. Explain the role of network protocols in facilitating IPC between processes in distributed systems.

Ans: Network protocols serve as standardized rules governing data transmission, defining formats, error handling, and communication procedures. In distributed

systems, protocols like TCP and UDP enable IPC by managing data flow, addressing, and ensuring reliable or lightweight communication between remote processes. These protocols handle packet transmission, delivery, and synchronization, facilitating efficient interprocess communication across networks.

Q4. How do remote procedure calls (RPCs) function as a form of IPC in distributed systems?

Ans: Remote Procedure Calls (RPCs) enable interprocess communication in distributed systems by allowing a process to execute procedures or functions on a remote system as if they were local. Through stubs and marshalling, RPC abstracts network communication complexities, providing a transparent method for invoking procedures across different machines, facilitating efficient communication and resource sharing between distributed processes.

Q5. What are the advantages and drawbacks of using message-oriented middleware (MOM) for IPC in distributed systems?

Ans: Advantages of Message-Oriented Middleware (MOM) in IPC include reliable asynchronous communication, fault tolerance, and decoupling of sender and receiver. However, drawbacks include potential message delays, increased complexity due to middleware setup, and performance overheads. MOM offers robustness but may introduce latency and administrative complexity, requiring careful consideration for use in distributed systems.

Q6. Discuss the significance of serialization and deserialization in IPC across distributed systems.

Ans: Serialization converts data structures into a format for transmission, aiding IPC in distributed systems by enabling data exchange between disparate platforms. Deserialization reconstructs the received serialized data back into its original format. These processes ensure uniform interpretation of data across varied systems, facilitating seamless communication. Serialization and deserialization play a crucial role in maintaining data integrity, enabling interoperability, and facilitating efficient communication between distributed processes.

Q7. What are sockets, and how do they facilitate communication in distributed systems?

Ans: Sockets represent endpoints for sending and receiving data across a network in

distributed systems. They provide an interface for processes to establish network connections, enabling communication using protocols like TCP or UDP. Sockets abstract network complexities by allowing processes to read from and write to network connections. They facilitate bidirectional data flow, supporting reliable, connection-oriented (TCP) or connectionless (UDP) communication between distributed entities, enabling robust network interactions.

Q3. Discuss the importance of socket programming APIs in developing distributed applications.

Ans: Socket programming APIs provide tools and functions for developers to create network-aware applications in distributed systems. They offer standardized methods for establishing, managing, and interacting with network connections. These APIs abstract underlying network complexities, enabling developers to focus on application logic. By offering a consistent interface for communication, socket programming APIs allow seamless integration of networking capabilities, enabling the development of robust, scalable, and interconnected distributed applications across diverse platforms and environments.

Q9. What is External Data Representation (EDR) in the context of distributed systems?

Ans: External Data Representation (EDR) is a standard format used for encoding and representing data structures in a platform-independent manner across distributed systems. It ensures uniform data representation by defining a common format for various data types, enabling seamless data exchange between heterogeneous systems. EDR facilitates interoperability by allowing different systems to understand and process data uniformly, regardless of their native data representations.

Q10. What is TCP stream communication in distributed systems?

Ans: TCP (Transmission Control Protocol) is a connection-oriented protocol that ensures reliable, ordered, and error-checked delivery of data between applications in a distributed system. It establishes a connection between two endpoints, guarantees data integrity, and manages packet retransmission in case of loss.

Q11. What are the key features of TCP in distributed systems?

Ans: The key features of TCP in distributed systems are:

- *Reliability:* TCP ensures data delivery without loss or corruption.

- *Ordered delivery:* It maintains the order of data packets sent between sender and receiver.
- *Error handling:* TCP detects and retransmits lost packets, ensuring reliability.
- *Flow control:* It regulates data flow between systems to prevent congestion and ensure efficient transmission.

Q12. What is UDP stream communication in distributed systems?

Ans: UDP (User Datagram Protocol) is a connectionless protocol in distributed systems that offer a lightweight, low-latency communication method. It sends data packets, called datagrams, without establishing a connection or ensuring delivery. UDP is used when speed and reduced overhead are prioritized over reliability.

Q13. What is External Data Representation (EDR) in distributed systems?

Ans: EDR is a standard for encoding and decoding data structures to facilitate data exchange between different computer architectures in a distributed system. It ensures that data can be transferred and understood uniformly across various systems, regardless of their native data representations.

Q14. What is the purpose of EDR in distributed systems?

Ans: The primary purpose of XDR is to enable interoperability between heterogeneous systems by defining a standard format for data representation. It allows different systems, regardless of their hardware or software platforms, to communicate and share data seamlessly.

Q15. How does EDR work in distributed systems?

Ans: EDR specifies a standard format for encoding basic data types (such as integers, floats, arrays, structures) into a platform-independent binary format. Before transmission, data is encoded into EDR format at the sender's end. At the receiver's end, the data is decoded from the EDR format into the native data representation.

Q16. What is CORBA in the context of EDR in distributed systems?

Ans: Common Object Request Broker Architecture (CORBA) is a middleware that enables communication between distributed objects in a network-agnostic manner. It defines a standard for EDR by allowing objects written in different programming languages to interact seamlessly across a network.

Q17. What does CORBA offer in terms of EDR in distributed systems?

Ans: CORBA provides a platform-independent and language-neutral approach for objects to communicate and exchange data using its Interface Definition Language (IDL). It facilitates transparent data representation between different systems.

Q18. Explain Java Object Serialization as an EDR approach in distributed systems.

Ans: Java Object Serialization is a mechanism in Java that allows objects to be converted into a byte stream for storage or transmission. In distributed systems, it serves as an EDR approach by enabling objects to be serialized (converted to a common byte format) and deserialized (reconstructed from bytes) across different Java-based systems.

Q19. Give any two advantages of Java Object Serialization in distributed systems?

Ans: *Platform independence:* Serialized objects can be transmitted between different Java-based systems regardless of underlying architectures.

Ease of use: Java provides built-in support for serialization, making it straightforward to exchange objects between distributed components.

Q20. How does XML function as an EDR approach in distributed systems?

Ans: XML (eXtensible Markup Language) is a text-based format for representing structured data. In distributed systems, XML serves as an EDR approach by providing a common language-independent format for encoding data that can be easily interpreted by diverse systems.

Q21. What are the characteristics of XML as an EDR approach in distributed systems?

Ans: The characteristics of XML as an EDR approach are:

- *Human-readable:* XML is human-readable and self-descriptive, making it understandable and usable by both humans and machines.
- *Extensibility:* XML allows for the creation of custom data structures, making it adaptable to various data representation needs in distributed systems.
- *Interoperability:* XML enables data interchange between heterogeneous systems, ensuring compatibility and ease of integration.

Q22. What is multicast communication in distributed systems?

Ans: Multicast communication is a networking technique that allows a sender to transmit a single message to multiple recipients simultaneously. It is a one-to-many

communication paradigm where data is efficiently distributed to a group of nodes that have expressed interest in receiving the information.

Q23. Give any two advantages of multicast communication in distributed systems?

Ans: The advantages of multicast communication are:

- *Reduced network traffic:* Multicast reduces bandwidth consumption by sending a single copy of data to multiple recipients, instead of individual unicast transmissions to each node.
- *Scalability:* It efficiently scales to accommodate varying group sizes without significant performance degradation, making it suitable for applications like video streaming, online gaming, and content distribution.

Long Answer Questions

Q1. What is inter-process communication? Explain its mechanism and benefits in distributed environments.

Ans: Refer Section 3.1

Q2. Explain the role and functionality of sockets in establishing communication between processes or applications across a network in distributed systems.

Ans: Refer Section 3.2.2

Q3. Explain the characteristics and advantages of TCP stream communication in distributed systems.

Ans: Refer Section 3.2.4

Q4. What are the primary characteristics and limitations of UDP stream communication in distributed systems?

Ans: Refer Section 3.2.3

Q5. Compare and contrast the use cases and scenarios where TCP stream communication is preferred over UDP stream communication in distributed systems.

Ans: Refer Section 3.2.4 and 3.2.3

Q6. What is external Data Representation (EDR)? Discuss the advantages of using EDR for data representation and communication in distributed systems.

Ans: Refer Section 3.3

Q7. What are the advantages of using CORBA for EDR in distributed systems?

Ans: Refer Section 3.4.1

- Q8. Describe the Java Object Serialization mechanism and its role in EDR for distributed systems.
- Ans: Refer Section 3.4.2
- Q9. Explain how XML is used for External Data Representation in distributed systems.
- Ans: Refer Section 3.4.3
- Q10. Explain the advantages and challenges associated with multicast communication in distributed systems.
- Ans: Refer Section 3.5
- Q11. Explain the benefits and challenges associated with network virtualization in distributed systems.
- Ans: Refer Section 3.6
- Q12. What are overlay networks, and how do they function in distributed systems?
- Ans: Refer Section 3.6.1

EXERCISE

1. Explain the impact of network latency and bandwidth on IPC in distributed systems.
2. Describe the concept of middleware in supporting IPC among various components of distributed systems.
3. Discuss the differences between TCP and UDP sockets in distributed systems.
4. Compare and contrast the characteristics, reliability, and use cases of Transmission Control Protocol (TCP) and User Datagram Protocol (UDP) sockets in distributed environments.
5. Discuss the reliability, ordered delivery, and connection-oriented nature of TCP streams emphasizing their significance in distributed environments.
6. Explain the concept of datagrams in UDP and how they contribute to stream communication in distributed systems.
7. How does EDR ensure data integrity and compatibility between different architectures in distributed systems?
8. Explain the key concepts and components of CORBA (Common Object Request Broker Architecture) in distributed systems.
9. Discuss the advantages and challenges associated with using Java Object Serialization for EDR in distributed systems.
10. Describe the key techniques and technologies used in network virtualization for distributed systems.